


ADVANCEDHMC.JL: A MODULAR IMPLEMENTATION OF STAN'S NO-U-TURN SAMPLER IN JULIA

Kai Xu¹, Hong Ge², Will Tebbutt², Mohamed Tarek³, Martin Trapp⁴ and Zoubin Ghahramani^{2,5}

¹University of Edinburgh ²University of Cambridge ³UNSW Canberra ⁴Graz University of Technology ⁵Uber AI Labs



Abstract

The No-U-Turn Sampler (NUTS) in **Stan** (Hoffman and Gelman, 2014; Carpenter et al., 2017) has demonstrated remarkable sampling robustness and efficiency in a wide range of Bayesian inference problems, due to the use of dynamic Hamiltonian trajectory, and a fine-tuned joint adaptation of step-size and mass matrix. Motivated by these successes, we present **AdvancedHMC.jl** (AHMC), a pure Julia implementation of **Stan**'s built-in NUTS algorithm and related adaptation methods. We hope **AdvancedHMC.jl** can help expose **Stan**'s NUTS to a wider range of users, e.g. those who want to write their models by hand, or using a different probabilistic programming language (e.g. **Turing**, **Soss**). In our package, NUTS is defined as a combination of individual components with abstractions partially inspired by (Betancourt, 2017).

Hamiltonian Monte Carlo Components

Hamiltonian Monte Carlo (HMC) simulates Hamiltonian dynamics to make proposals for a Markov chain (Neal et al., 2011). **AdvancedHMC.jl** supports various HMC samplers below.

$$(\text{StaticHMC} \cup \text{DynamicHMC}) \times \text{Adaptor.}$$

Here **StaticHMC** are HMC with fixed-length trajectories and **DynamicHMC** are HMC with adaptive trajectory length which can be created by composing NUTS components as follows

$$\text{Metric} \times \text{Integrator} \times \text{TrajectorySampler} \times \text{TerminationCriterion},$$

where

$$\begin{aligned} \text{Metric} &= \{\text{UnitEuclidean}, \text{DiagEuclidean}, \text{DenseEuclidean}\} \\ \text{Integrator} &= \{\text{Leapfrog}\} \\ \text{TrajectorySampler} &= \{\text{Slice}, \text{Multinomial}\} \\ \text{TerminationCriterion} &= \{\text{ClassicNoUTurn}, \text{GeneralisedNoUTurn}\} \end{aligned}$$

Adaptor can be composed from base adaptors

$$\text{BaseAdaptor} \in \{\text{Preconditioner}, \text{NesterovDualAveraging}\}.$$

Note 1: **Preconditioner** behaves differently based on the choice of metric spaces.

Note 2: **StanHMCAdaptor**, a specific composition of base adaptors that is equivalent to **Stan**'s windowed adaptor, is provided. This adaptor has been proved to be robust in practice.

Benchmark Models

We use five models from **MCMCBenchmarks.jl** to compare NUTS between **AdvancedHMC.jl** and **Stan**.

Gaussian Model (Gaussian) is a simple two parameter Gaussian distribution.

$$\mu \sim \mathcal{N}(0, 1), \quad \sigma \sim \text{Truncated}(\text{Cauchy}(0, 5), 0, \infty), \quad y_n \sim \mathcal{N}(\mu, \sigma) \quad (n = 1, \dots, N)$$

Signal Detection Model (SDT) is a model used in psychophysics and signal processing, which decomposes performance in terms of discriminability and bias.

$$d \sim \mathcal{N}(0, \frac{1}{\sqrt{2}}), \quad c \sim \mathcal{N}(0, \frac{1}{\sqrt{2}}), \quad x \sim \text{SDT}(d, c)$$

Linear Regression Model (LR) is a linear regression with truncated Cauchy prior on the weights.

$$B_d \sim \mathcal{N}(0, 10), \quad \sigma \sim \text{Truncated}(\text{Cauchy}(0, 5), 0, \infty), \quad y_n \sim \mathcal{N}(\mu_n, \sigma),$$

where $\mu = B_0 + B^T X$, $d = 1, \dots, D$ and $n = 1, \dots, N$.

Hierarchical Poisson Regression (HPR)

$$a_0 \sim \mathcal{N}(0, 10), \quad a_1 \sim \mathcal{N}(0, 1), \quad b_\sigma \sim \text{Truncated}(\text{Cauchy}(0, 1), 0, \infty), \quad b_d \sim \mathcal{N}(0, b_\sigma), \quad y_n \sim \text{Poi}(\log \lambda_n),$$

where $\log \lambda_n = a_0 + b_{z_n} + a_1 x_n$, $d = 1, \dots, N_b$ and $n = 1, \dots, N$.

Linear Ballistic Accumulator (LBA) is a cognitive model of perception and simple decision making.

$$\tau \sim \text{Truncated}(\mathcal{N}(0.4, 0.1), 0, mn), \quad A \sim \text{Truncated}(\mathcal{N}(0.8, 0.4), 0, \infty),$$

$$k \sim \text{Truncated}(\mathcal{N}(0.2, 0.3), 0, \infty), \quad \nu_d \sim \text{Truncated}(\mathcal{N}(0, 3), 0, \infty), \quad x_n \sim \text{LBA}(\nu, \tau, A, k)$$

where $mn = \min_i x_{i,2}$, $d = 1, \dots, N_c$ and $n = 1, \dots, N$.

Example Code of Building Stan's NUTS using AHMC

```
1 using AdvancedHMC
2 n_samples, n_adapts, target = 10_000, 2_000, 0.8 # set up sampling parameters
3 q_init = randn(D) # draw a random starting point
4 ### Building up NUTS
5 metric = DiagEuclideanMetric(D) # diagonal Euclidean metric space
6 h = Hamiltonian(metric, logdensity_f, grad_f) # Hamiltonian on the target distribution
7 eps_init = find_good_eps(h, q_init) # initial step size
8 int = Leapfrog(eps_init) # leapfrog integrator
9 traj = NUTS{Multinomial, GeneralisedNoUTurn}(int) # multinomial sampling with generalised no U-turn
10 adaptor = StanHMCAdaptor( # Stan's windowed adaptor
11     n_adapts, Preconditioner(metric), NesterovDualAveraging(target, eps_init)
12 )
13 samples, stats = sample(h, traj, q_init, n_samples, adaptor, n_adapts) # draw samples
```

Sampling Efficiency: Stan's NUTS v.s. AHMC

To compare the sampling efficiency between **Stan** and **AHMC**, we run multiple runs of NUTS with target acceptance rate 0.8 for 2,000 runs with 1,000 adaptation steps, where the warm-up samples dropped. Below are figures of distributions of step size and tree depth, and the mean effective sample size (ESS) for different variables.

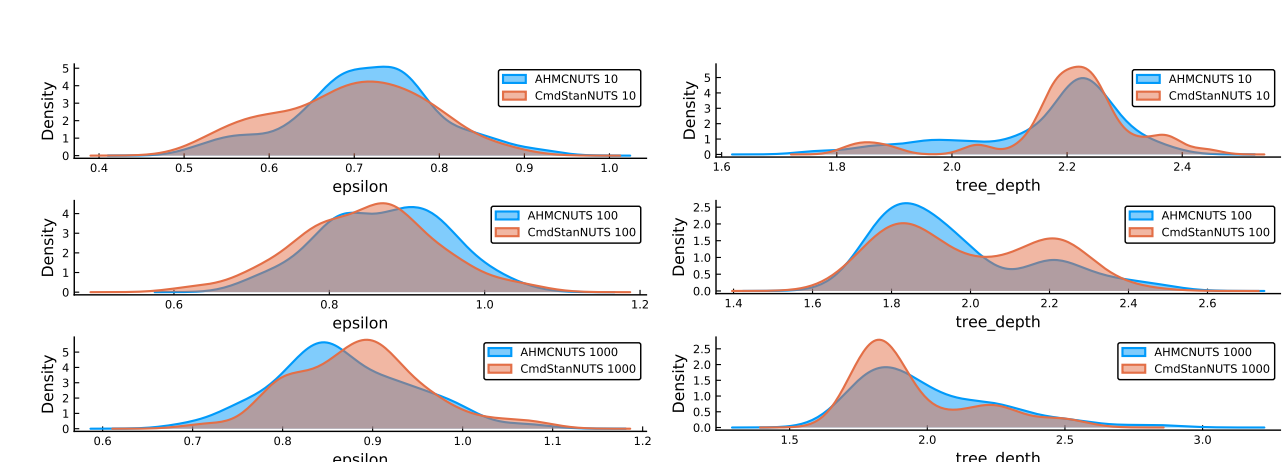


Fig. 1: Gaussian (50 runs); left to right: step size, tree depth, ESS

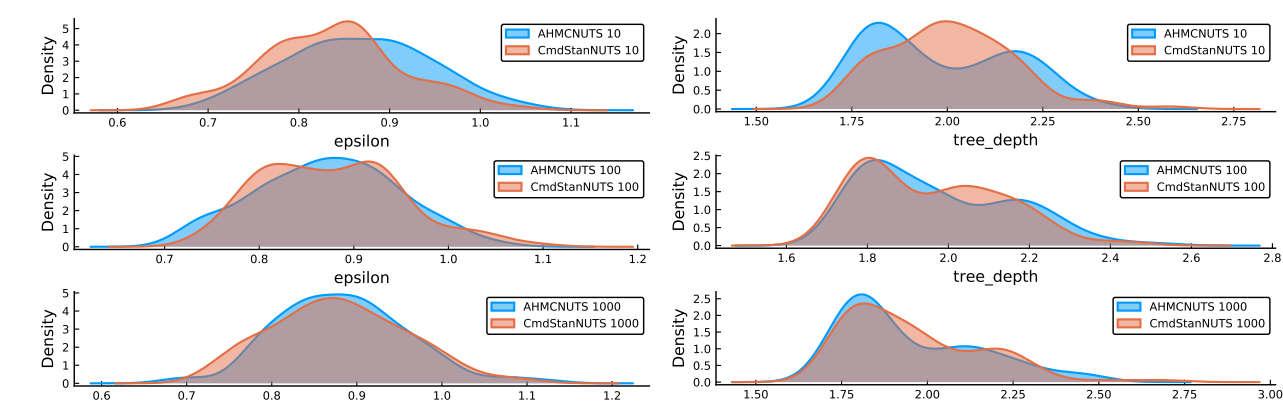


Fig. 2: SDT (100 runs); left to right: step size, tree depth, ESS

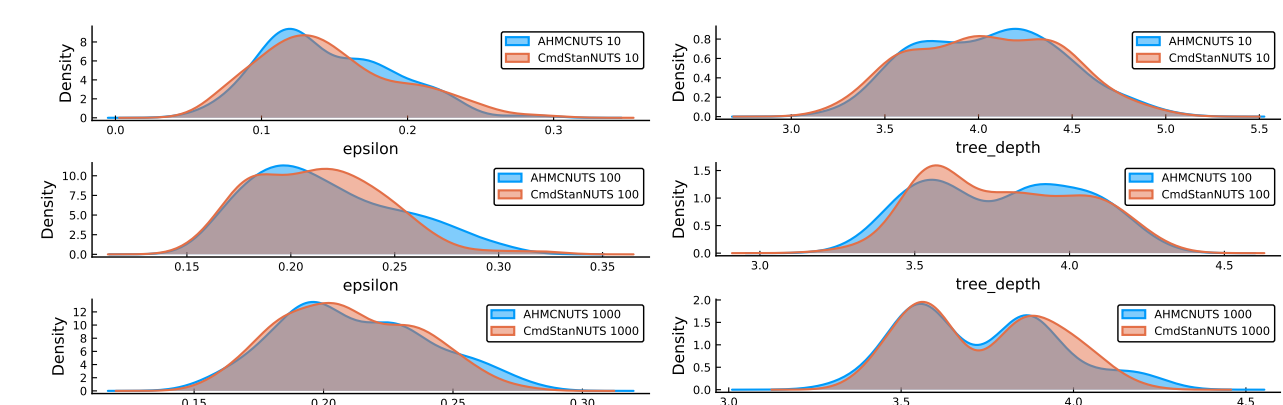


Fig. 3: LR (50 runs); left to right: step size, tree depth, ESS

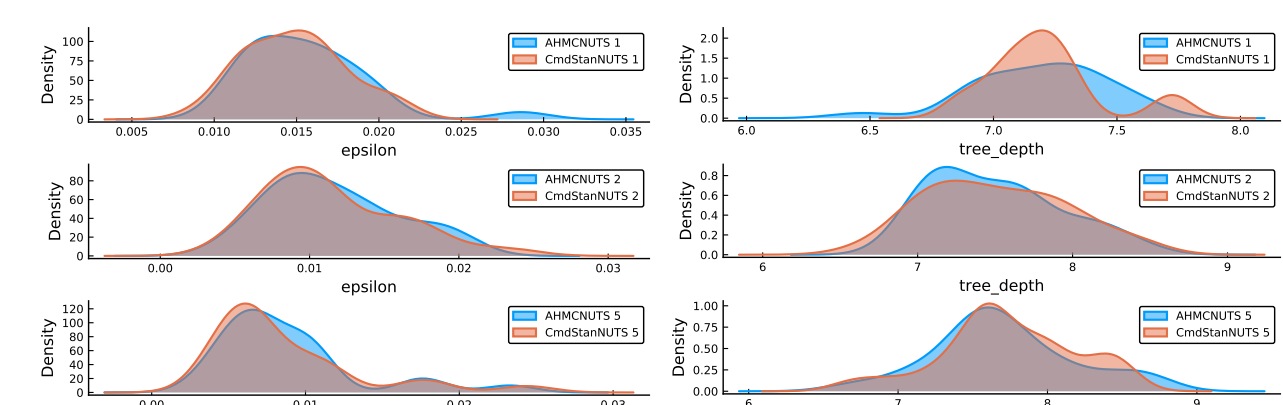


Fig. 4: HPR (25 runs); left to right: step size, tree depth, ESS (of some variables)

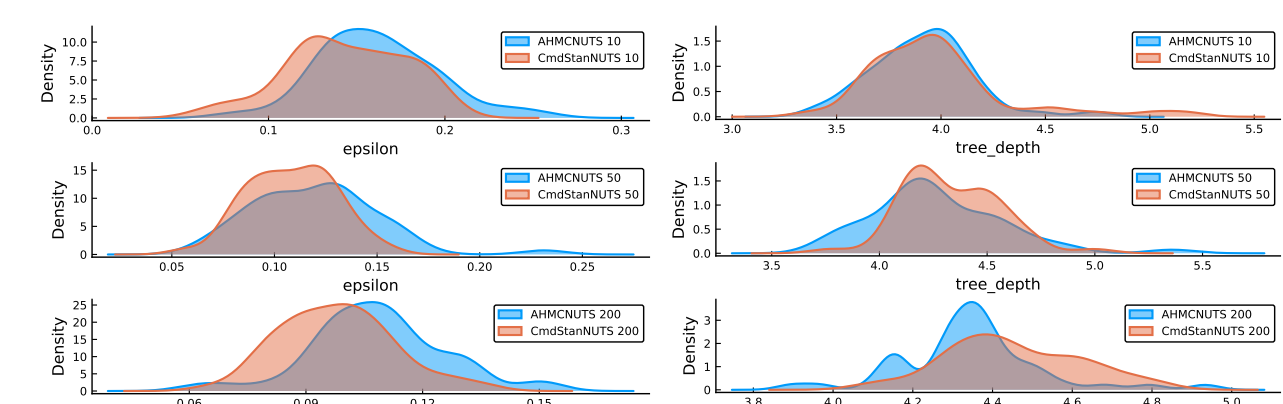


Fig. 5: LBA (50 runs); left to right: step size, tree depth, ESS (of some variables)

Computational Efficiency: Stan v.s. Turing

Turing.jl is a probabilistic programming language (PPL) in Julia that uses **AdvancedHMC.jl** as its HMC backend. All the benchmark models are written in **Turing** and **AdvancedHMC.jl** is called by **Turing.jl** to run the NUTS. Below is an example of running NUTS on the LR model using **Turing**.

```
1 @model LR(x, y, Nd, Nc) = begin
2     B ~ MvNormal(zeros(Nc), 10)
3     B0 ~ Normal(0, 10)
4     sigma ~ Truncated(Cauchy(0, 5), 0, Inf)
5     mu = B0 .+ x * B
6     y ~ MvNormal(mu, sigma)
7 end
8 x, y, Nd, Nc = ... # load data
9 chain = sample(LR(x, y, Nd, Nc), NUTS(2_000, 1_000, 0.8))
```

The time to run the five benchmark models in **Stan** and **Turing** are reported in the table below.

	Gaussian ²		SDT ³		LR ²		HPR ¹		LBA ²	
	N	seconds	N	seconds	N	seconds	N	seconds	N	seconds
Stan	10	0.8039	10	0.7759	10	0.8669	10	2.4870	10	1.9179
AHMC	10	0.3361	10	0.3285	10	1.1356	10	19.4587	10	2.6906
Stan	100	0.7561	100	0.7261	100	0.9824	20	3.5025	50	7.8471
AHMC	100	0.3303	100	0.3201	100	1.3202	20	28.2982	50	11.0270
Stan	1000	0.7614	1000	0.7089	1000	2.2600	50	5.8954	200	31.3762
AHMC	1000	0.5081	1000	0.3179	1000	3.8326	50	40.0322	200	33.6125

Table 1: Time comparisons between **Stan** and **Turing** (AHMC) for five models using ¹ 25 runs, ² 50 runs or ³ 100 runs.

Easy Integration of Other Julia Packages

Bijectors.jl is used inside **Turing.jl** to do automatic transformations of constrained variables to run HMC. E.g. a random variable from $\text{Truncated}(\text{Cauchy}(0, 5), 0, \infty)$ is constrained to be positive and will be transformed to the real space by the log function automatically.

CuArrays.jl could be used with **AdvancedHMC.jl** to run NUTS on GPUs. In order to run NUTS using CUDA, one only needs to change Line 3 of the demo code from `q_init = randn(D)` to `q_init = CuArray(randn(D))`, assuming `logdensity_f` and `grad_f` in Line 6 are GPU friendly; if it is written in pure Julia, it probably supports GPUs acceleration automatically.

How does it work? All arrays in **AdvancedHMC.jl** are abstractly typed, meaning that the concrete type is deduced at compile time from `q_init`. That is to say, if `q_init` is on the GPU i.e. is a `CuArray`, all the internal arrays in the NUTS will be too.

SoSS.jl is another PPL in Julia that uses **AdvancedHMC.jl** as its backend. It is easy for PPLs in Julia with different domain specific languages (DSLs) to use the HMC implementation in **AdvancedHMC.jl**.

DifferentialEquations.jl is the state-of-the-art numerical differential equations solver package, implemented in pure Julia. As such, its solvers can be employed in Turing models, thus enabling **AHMC** to perform Bayesian inference in the parameters of differential equation models.

Flux.jl is a deep learning packages in Julia. Neural models defined by **Flux.jl** can be directly used in Turing models. E.g. one can implement a Bayesian neural network in Turing by defining priors on the weights of a Flux-based neural network, and NUTS can be used to draw samples of the weights.

Acknowledgements

We would like to thank the developers of **MCMCBenchmarks.jl**, Rob Goedman and Christopher Fisher. With this package we ran all benchmarks and generated all plots presented in this poster.

References

- Betancourt, M. (2017). A conceptual introduction to Hamiltonian Monte Carlo. *arXiv preprint arXiv:1701.02434*.
- Carpenter, B., Gelman, A., Hoffman, M. D., Lee, D., Goodrich, B., Betancourt, M., ... Riddell, A. (2017). Stan: a probabilistic programming language. *Journal of statistical software*, 76(1).
- Hoffman, M. D. & Gelman, A. (2014). The No-U-Turn sampler: adaptively setting path lengths in Hamiltonian Monte Carlo. *Journal of Machine Learning Research*, 15(1), 1593-1623.
- Neal, R. M. et al. (2011). MCMC using Hamiltonian dynamics. *Handbook of Markov chain Monte Carlo*, 2(11), 2.